

▼ 1. Make neccesary imports and installs

To use the models and different plots and reports, we make the neccesary imports and insstall the libraries that are not installed by default

We also use a seed value which makes the code reproducible.

```
!pip install ydata-profiling  
!pip install --upgrade ydata-profiling
```

```
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-pa  
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/d  
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/d  
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.12/dist-pa
```

```
Requirement already satisfied: pydantic-core==2.33.2 in /usr/local/lib/python3.11/site-packages
Requirement already satisfied: typing-extensions>=4.12.2 in /usr/local/lib/python3.11/site-packages
Requirement already satisfied: typing-inspection>=0.4.0 in /usr/local/lib/python3.11/site-packages
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.11/site-packages
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: attrs>=19.3.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: networkx>=2.4 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: puremagic in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from ydata_profiling import ProfileReport

from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OrdinalEncoder, OneHotEncoder, LabelEncoder
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import StratifiedKFold

from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.model_selection import GridSearchCV

np.random.seed(0)
```

[Upgrade to ydata-sdk](#)

Improve your data and profiling with ydata-sdk, featuring data quality scoring, redundancy detection, outlier identification, text validation, and synthetic data generation.

2. Import the data and do the initial analysis

We import the data and then convert it to a dataframe that can easily be used by python to perform different operations.

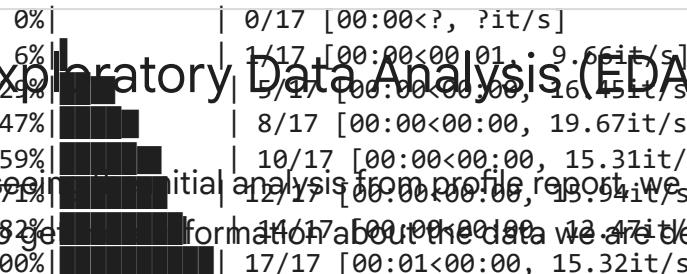
We use profile report which can be used to do some basic EDA on the data like checking missing values, duplicates, basic column statistics and correlations.

```
df = pd.read_csv("/content/drive/MyDrive/project/train.csv")
df = df.drop('id', axis=1)
```

```
profile = ProfileReport(df, title="Pandas Profiling Report")
profile.to_file("your_report.html")
profile
```


Summarize dataset: 100%

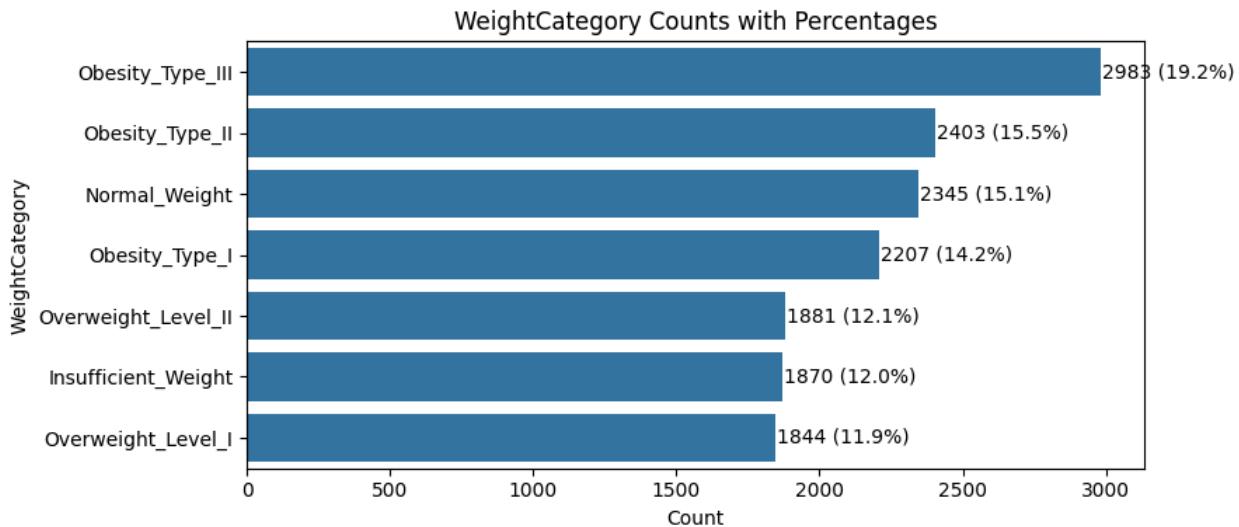
90/90 [00:33<00:00, 2.37it/s, Completed]



3. Exploratory Data Analysis (EDA)

After seeing initial analysis from profile report, we do some additional analysis on the data to get information about the data we are dealing with.

```
plt.figure(figsize=(9,4))
ax = sns.countplot(y='WeightCategory', data=df, order=df['WeightCategory'].value_
total = len(df)
for p in ax.patches:
    count = int(p.get_width())
    pct = 100 * count / total
    ax.text(p.get_width()+5, p.get_y()+p.get_height()/2, f'{count} ({pct:.1f}%)'
            va='center')
ax.set_title('WeightCategory Counts with Percentages')
ax.set_xlabel('Count')
plt.tight_layout()
plt.show()
```



Duplicate rows 0

Duplicate rows (%) 0.0%

Total size in memory 2.0 MiB

Above we checked if our prediction column in the test dataset is balanced or not. While all the categories are not equally distributed, the distribution is not completely biased towards a particular class either.

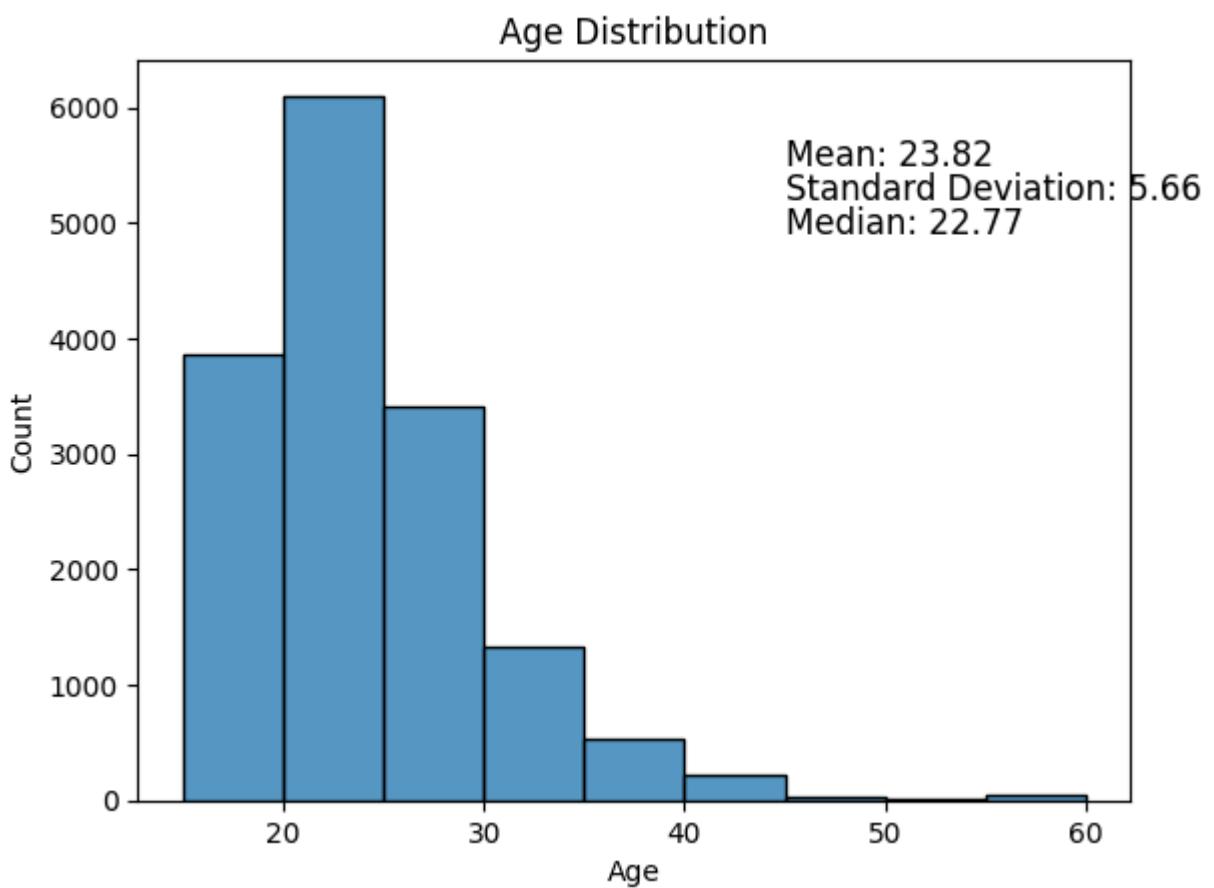
Obesity_Type_III has a lot more entries and we need to take care about that while creating our training data.

Variabiles

```
bin_edges = [15, 20, 25, 30, 35, 40, 45, 50, 55, 60]
sns.histplot(df['Age'], bins=bin_edges, kde=False)
plt.title("Age Distribution")

mean_age = df['Age'].mean()
std_age = df['Age'].std()
median_age = df['Age'].median()

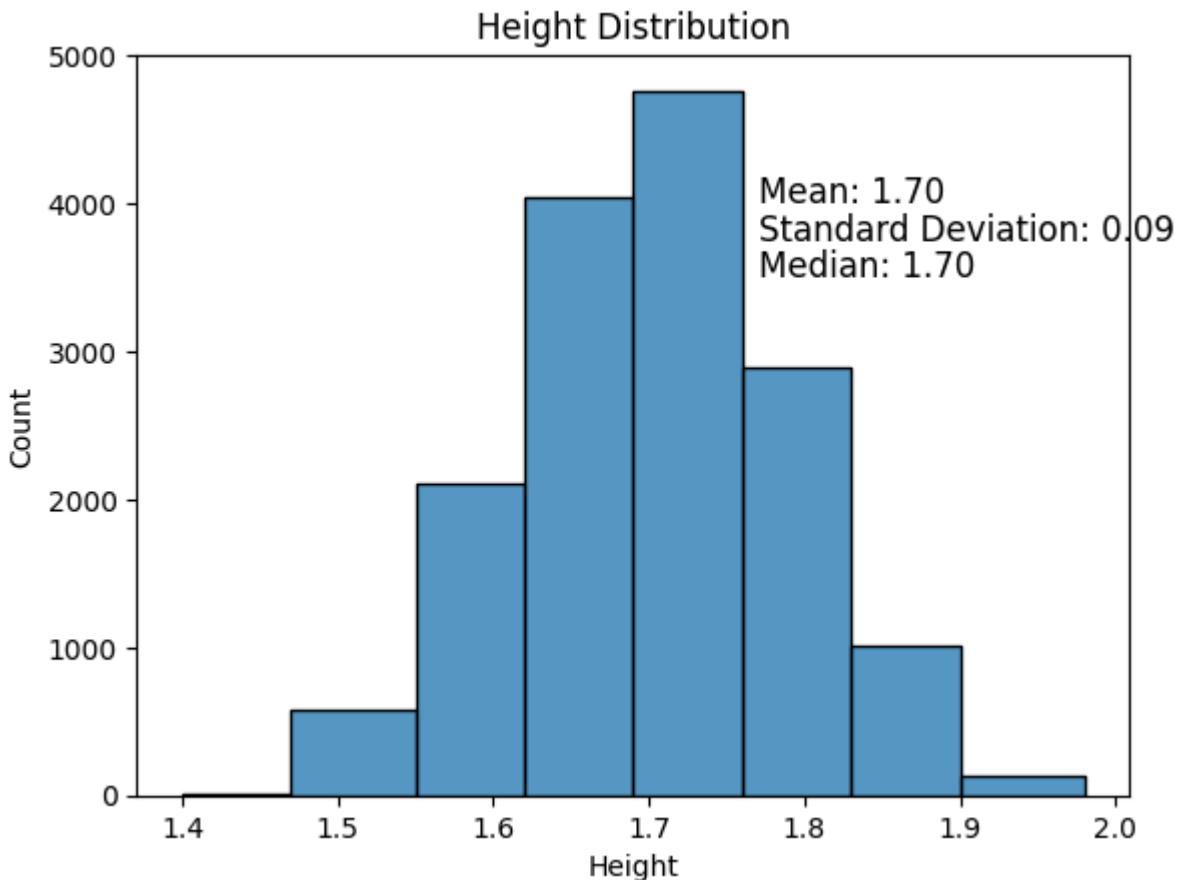
plt.text(45, 5500, f'Mean: {mean_age:.2f}', fontsize=12)
plt.text(45, 5200, f'Standard Deviation: {std_age:.2f}', fontsize=12)
plt.text(45, 4900, f'Median: {median_age:.2f}', fontsize=12)
plt.show()
```



```
# Histogram for Height
bin_edges = [1.4, 1.47, 1.55, 1.62, 1.69, 1.76, 1.83, 1.9, 1.98]
sns.histplot(df['Height'], bins=bin_edges, kde=False)
plt.title("Height Distribution")

mean_height = df['Height'].mean()
std_height = df['Height'].std()
median_height = df['Height'].median()

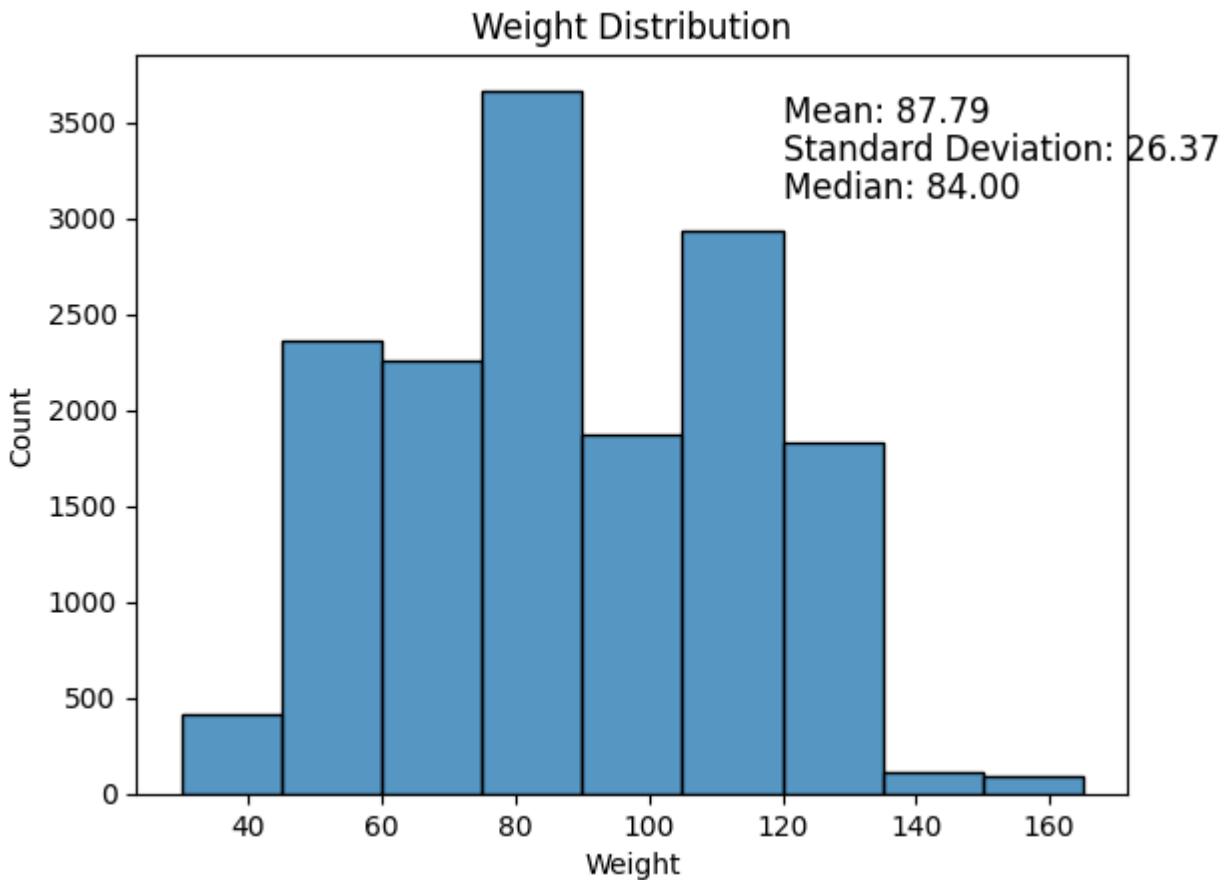
plt.text(1.77, 4000, f'Mean: {mean_height:.2f}', fontsize=12)
plt.text(1.77, 3750, f'Standard Deviation: {std_height:.2f}', fontsize=12)
plt.text(1.77, 3500, f'Median: {median_height:.2f}', fontsize=12)
plt.show()
```



```
bin_edges = [30, 45, 60, 75, 90, 105, 120, 135, 150, 165]
sns.histplot(df['Weight'], bins=bin_edges, kde=False)
plt.title("Weight Distribution")

mean_weight = df['Weight'].mean()
std_weight = df['Weight'].std()
median_weight = df['Weight'].median()

plt.text(120, 3500, f'Mean: {mean_weight:.2f}', fontsize=12)
plt.text(120, 3300, f'Standard Deviation: {std_weight:.2f}', fontsize=12)
plt.text(120, 3100, f'Median: {median_weight:.2f}', fontsize=12)
plt.show()
```



We plot the columns age, height and weight and check the spread of the data in these columns. We use histogram with bins of equal intervals to easily make sense of which bin has more frequency.

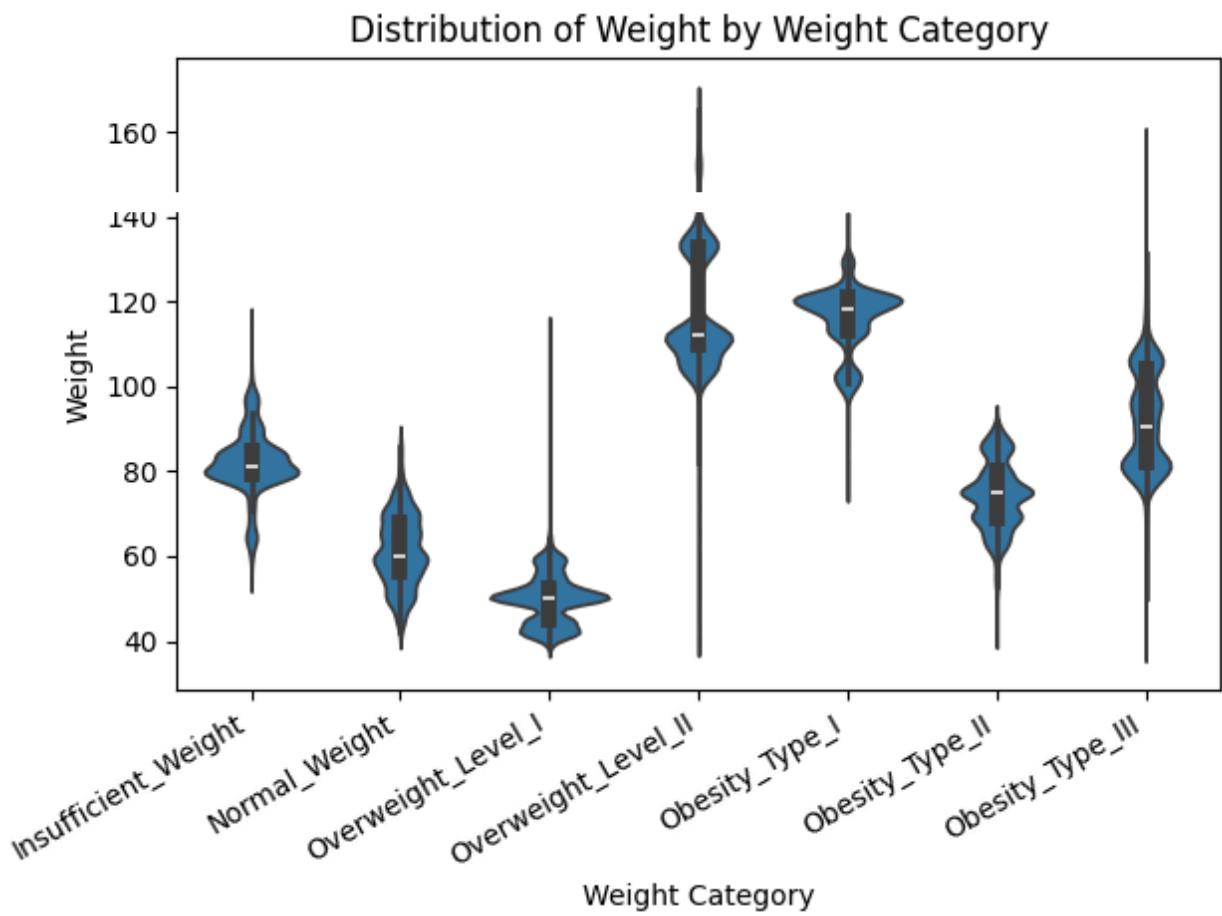
The graph for age is an off-center bell curve with bias towards 20-30 years age group, which implies the dataset has more people in their 20s. The standard deviation is also around 5 which is not that much considering the age range. So instead of using bins, we are better off giving this data in form of continuous values since bins would create a bias.

Height column seems to be more of a traditional bell curve with the average around 1.7. The median is also 1.7 with standard deviation very low which means most people have their heights around mean. So this column does not seem to be a big differentiator.

The weight column is where it gets interesting. The bins form a pretty flattened bell curve with values having a lot less difference in frequencies between bins. The high standard deviation indicating a wider range and greater variability in weight among the individuals in the dataset. Let us compare it with the resultant weight category column since both of them seems to have a similar split.

```
rev_map = {
    0: 'Insufficient_Weight',
    1: 'Normal_Weight',
    2: 'Overweight_Level_I',
    3: 'Overweight_Level_II',
    4: 'Obesity_Type_I',
```

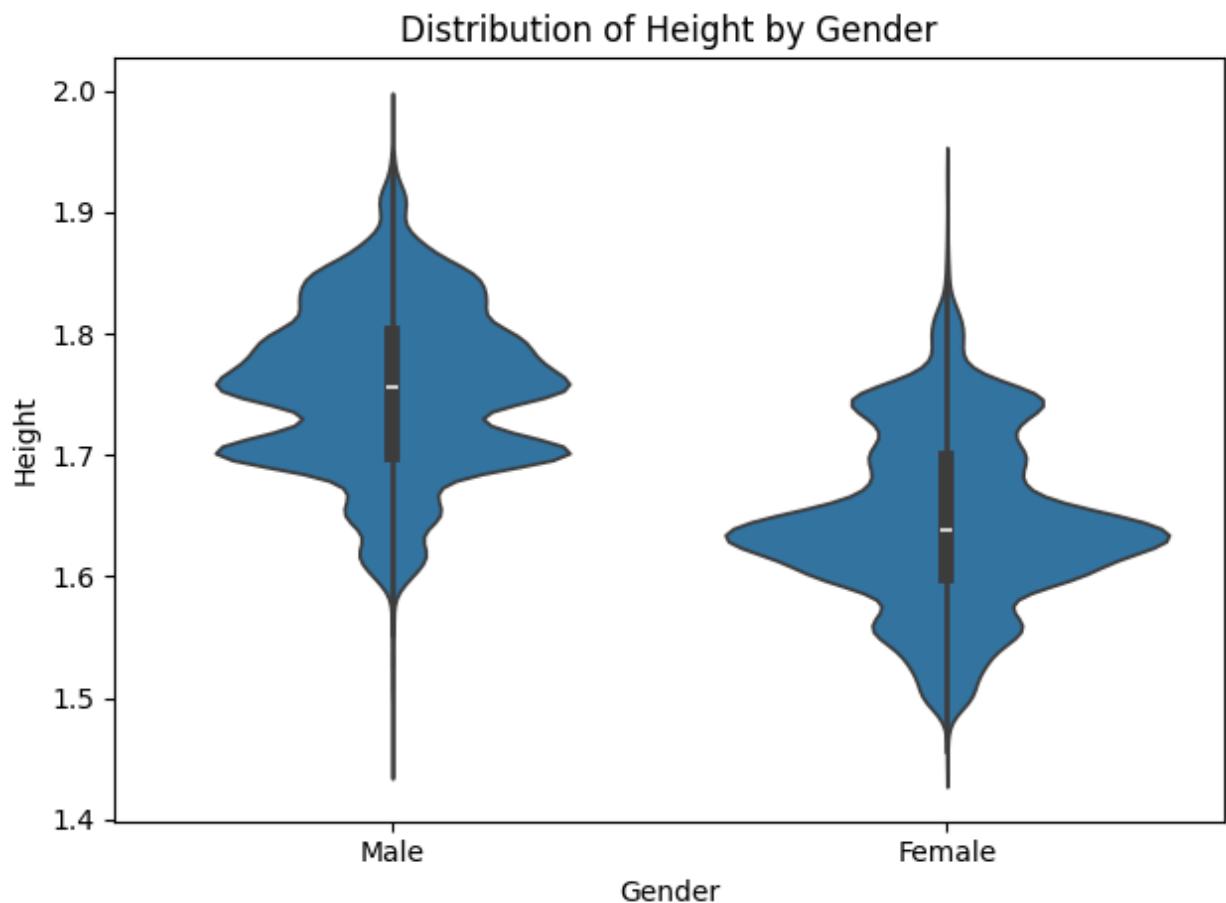
```
5: 'Obesity_Type_II',
6: 'Obesity_Type_III'
}
sns.violinplot(x='WeightCategory', y='Weight', data=df)
plt.title('Distribution of Weight by Weight Category')
plt.xlabel('Weight Category')
plt.ylabel('Weight')
plt.xticks(ticks=list(rev_map.keys()), labels=list(rev_map.values()), rotation=30)
plt.tight_layout()
plt.show()
```



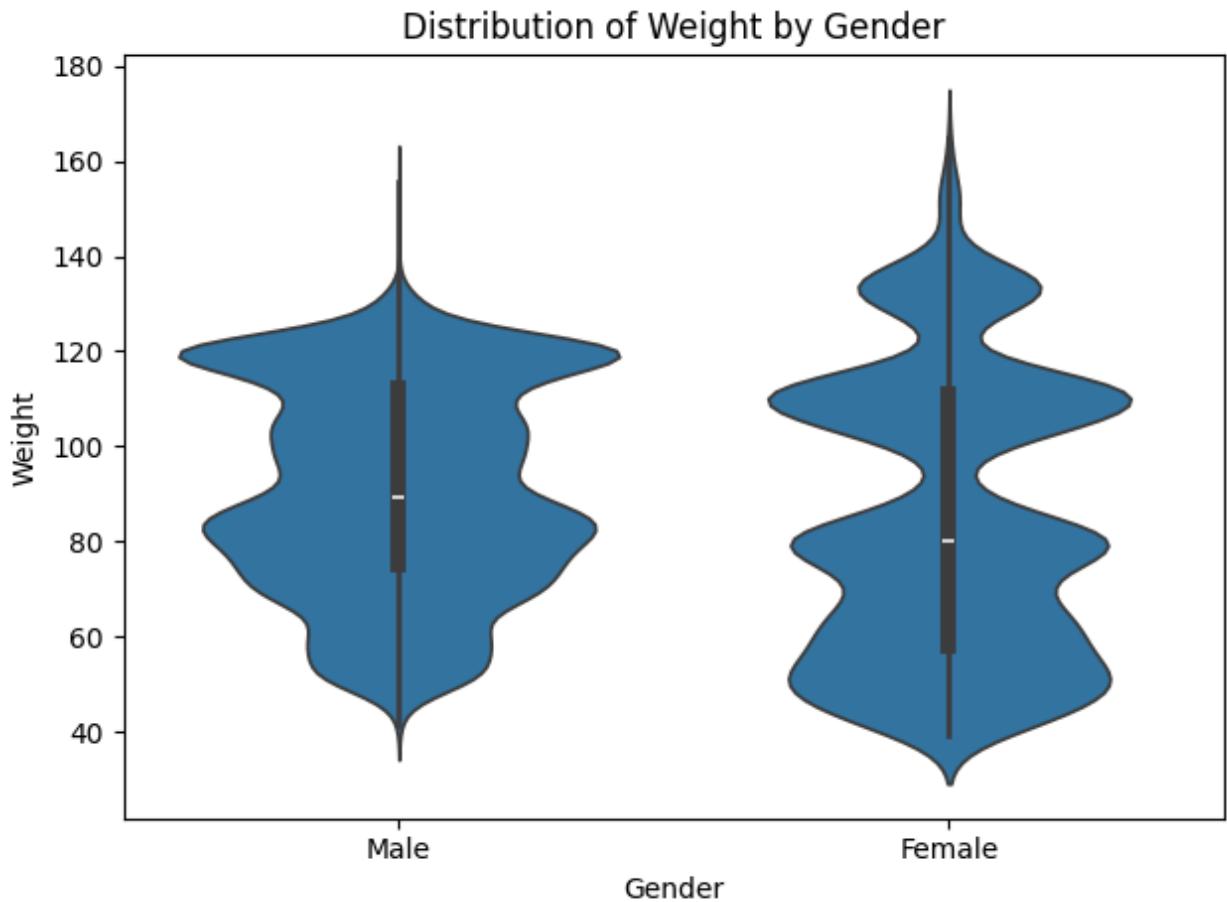
The violin plot shows a relationship between weight and weight categories. With each weight category, the average weight seems to go higher with more datapoints around the average for each category.

Now let us try visualising how gender is affecting our data.

```
sns.violinplot(x='Gender', y='Height', data=df)
plt.title('Distribution of Height by Gender')
plt.xlabel('Gender')
plt.ylabel('Height')
plt.tight_layout()
plt.show()
```



```
sns.violinplot(x='Gender', y='Weight', data=df)
plt.title('Distribution of Weight by Gender')
plt.xlabel('Gender')
plt.ylabel('Weight')
plt.tight_layout()
plt.show()
```



We try to plot the graph for height and weight by gender. We can notice that the height is varying by gender as expected. Males in average have higher height than females and the shape of the violin suggests that the spread is around average as expected.

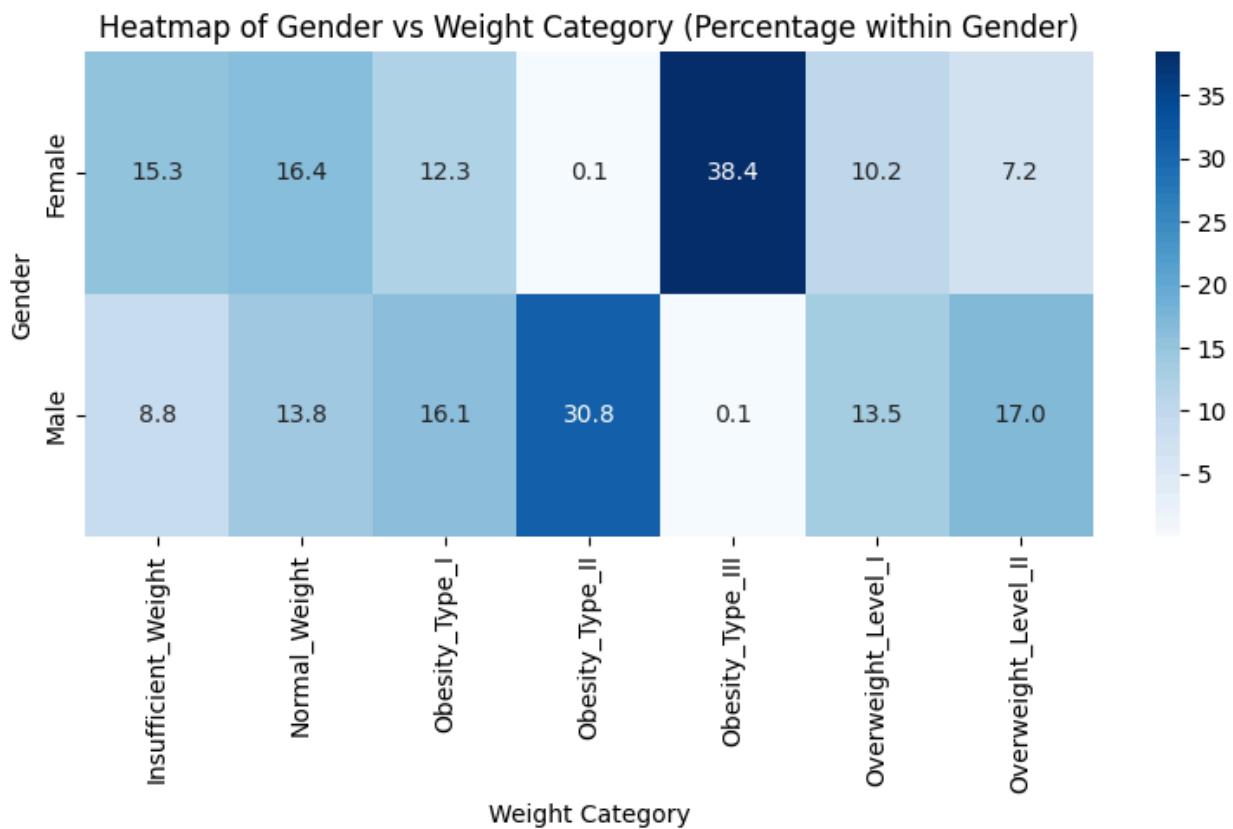
For weight we can see that while the average is pretty similar the data for the females have a lot of variations and the data tends to skew in both the extremes in case of females. So we should try and do more analysis on relation between weights and the gender.

```
crosstab_family = pd.crosstab(df['Gender'], df['WeightCategory'])
crosstab_family = crosstab_family.rename(columns=rev_map)
crosstab_family_percentage = crosstab_family.div(crosstab_family.sum(axis=1), axis=0)

plt.figure(figsize=(8, 5))
sns.heatmap(crosstab_family_percentage, annot=True, fmt=' .1f', cmap='Blues')

plt.title('Heatmap of Gender vs Weight Category (Percentage within Gender)')
plt.xlabel('Weight Category')
plt.ylabel('Gender')

plt.tight_layout()
plt.show()
```



From the heatmap we can see that in case of the female dataset, the dataset seems to skew more towards both the extremes. Females tend to have more insufficient weights in this dataset than men and a lot more obesity type 3 datapoints than males.

Males tend to be more type 2 obese and females tend to be more type 3 obese in this dataset. This would mean Gender would also be a good feature.

Now we go ahead and do analysis on more columns.

```

crosstab = pd.crosstab(df['FAVC'], df['WeightCategory'])
crosstab = crosstab.rename(columns=rev_map)
crosstab_percentage = crosstab.div(crosstab.sum(axis=1), axis=0) * 100

plt.figure(figsize=(8, 5))
sns.heatmap(crosstab_percentage, annot=True, fmt=' .1f', cmap='Blues')

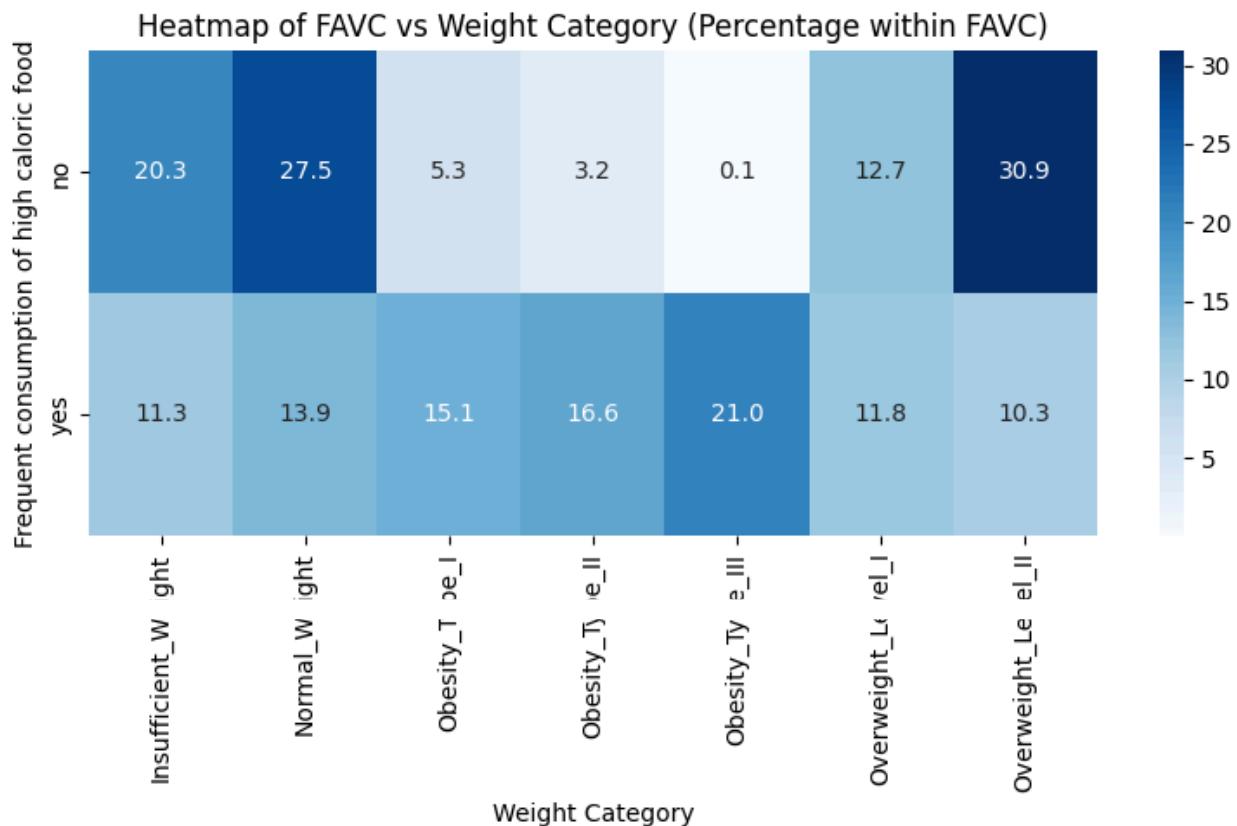
plt.title('Heatmap of FAVC vs Weight Category (Percentage within FAVC)')
plt.xlabel('Weight Category')
plt.ylabel('Frequent consumption of high caloric food')

plt.tight_layout()
plt.show()

favc_counts = df['FAVC'].value_counts()
favc_percentages = df['FAVC'].value_counts(normalize=True) * 100

```

```
favc_summary = pd.DataFrame({'Count': favc_counts, 'Percentage (%)': favc_percent})
print("Unique value counts and percentages for FAVC:")
display(favc_summary)
```



Unique value counts and percentages for FAVC:

	Count	Percentage (%)	Actions
FAVC			
yes	14184	91.315264	
no	1349	8.684736	

Next steps: [Generate code with favc_summary](#) [New interactive sheet](#)

We can see that the FAVC column (Frequent consumption of high caloric food) is highly imbalanced, 91% of the rows have the value as yes.

But important thing to notice is that when the FAVC feature is no, then the weight category is rarely in the obese category. So while this column is imbalanced, it provides vital information that helps us separate between obese and not obese.

```
crosstab_family = pd.crosstab(df['family_history_with_overweight'], df['WeightCategory'])
crosstab_family = crosstab_family.rename(columns=rev_map)
crosstab_family_percentage = crosstab_family.div(crosstab_family.sum(axis=1), axis=0)

plt.figure(figsize=(8, 5))
```

```

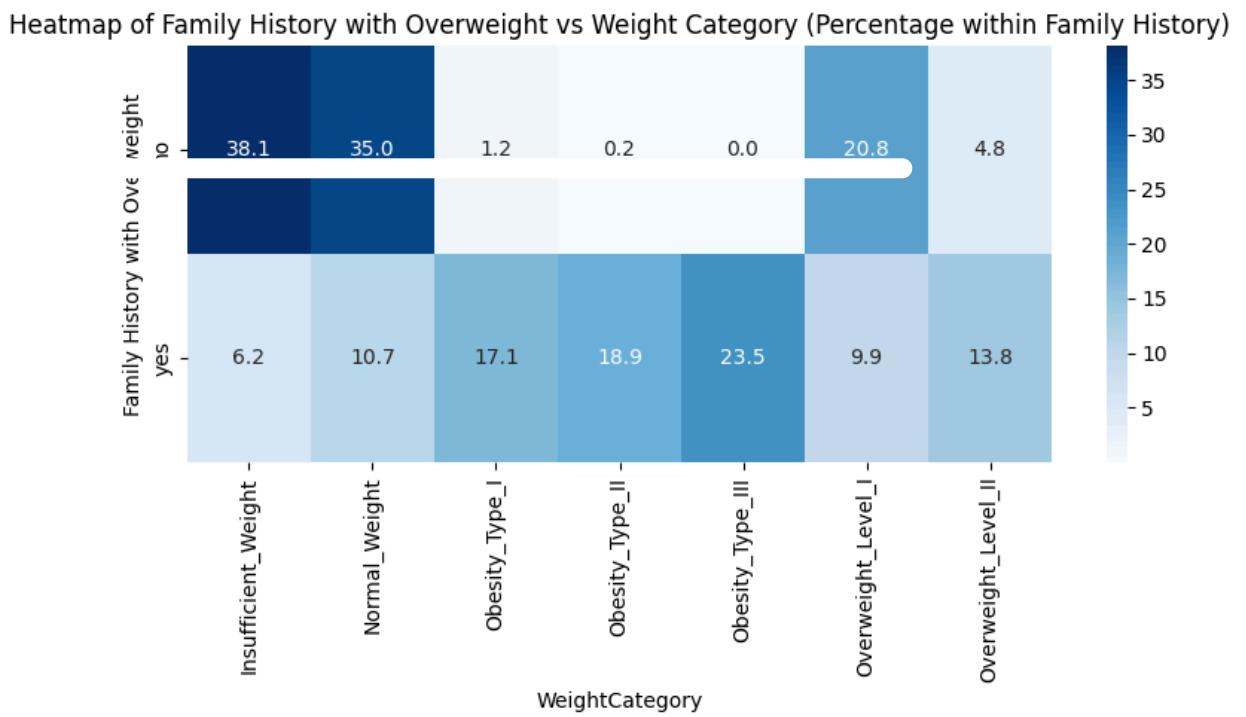
sns.heatmap(crosstab_family_percentage, annot=True, fmt='.1f', cmap='Blues')

plt.title('Heatmap of Family History with Overweight vs Weight Category (Percentage within Family History)')
plt.ylabel('Family History with Overweight')

plt.tight_layout()
plt.show()

fh_counts = df['family_history_with_overweight'].value_counts()
fh_percentages = df['family_history_with_overweight'].value_counts(normalize=True)
fh_summary = pd.DataFrame({'Count': fh_counts, 'Percentage (%)': fh_percentages})
print("Unique value counts and percentages for Family History:")
display(fh_summary)

```



Unique value counts and percentages for Family History:

Count	Percentage (%)
-------	----------------

family_history_with_overweight

yes	12696	81.73566
no	2837	18.26434

Next steps: [Generate code with fh_summary](#) [New interactive sheet](#)

Just like FAVC, family history column is also skewed but it also has a strong correlation with weight category. People who lack family history about obesity, also happen to belong to

mostly insufficient, normal and overweight category.

On the other hand people with family history of obesity tend to be obese themselves.

4. Preparing data

Now we prepare the data by converting categorical columns to numerical columns and using standard scaler to scale down different columns so all of them belong to the same scale. Then we separate the data into features, output, train and test dataset and so on. We make a reusable function so that the same operations can also be done on the test dataset.

```
weight_map = {
    'Insufficient_Weight': 0,
    'Normal_Weight': 1,
    'Overweight_Level_I': 2,
    'Overweight_Level_II': 3,
    'Obesity_Type_I': 4,
    'Obesity_Type_II': 5,
    'Obesity_Type_III': 6
}
df['WeightCategory'] = df['WeightCategory'].map(weight_map)
```

```
print(df.shape)
df.columns

(15533, 17)
Index(['Gender', 'Age', 'Height', 'Weight', 'family_history_with_overweight',
       'FAVC', 'FCVC', 'NCP', 'CAEC', 'SMOKE', 'CH20', 'SCC', 'FAF', 'TUE',
       'CALC', 'MTRANS', 'WeightCategory'],
      dtype='object')
```

```
nutrition_features = df[['FCVC', 'CH20', 'FAVC']]
nutrition_features['FAVC'] = nutrition_features['FAVC'].map({'yes': -1, 'no': 0})
scaler = StandardScaler()
nutrition_scaled = scaler.fit_transform(nutrition_features)
pca = PCA(n_components=1)
nutrition_score_pca = pca.fit_transform(nutrition_scaled)
print(pd.DataFrame({'Feature': ['FCVC', 'CH20', 'FAVC_numeric'], 'Weight': pca.
```

	Feature	Weight
0	FCVC	0.659168
1	CH20	0.695757
2	FAVC_numeric	-0.285340

```
/tmp/ipython-input-2669758093.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/>

```
nutrition_features['FAVC'] = nutrition_features['FAVC'].map({'yes': -1, 'no': 0})
```

Above we tried to find the best values to multiply the columns to create a combined health score using PCA.

```
def preprocess_data(df):

    #Creating Health score based on PCA weights
    df['HealthyEatingScore_PCA'] = df['FCVC']*0.659168 + df['CH20']*0.695757 + d

    #Creating X and y from the dataframe
    y = None
    if 'WeightCategory' in df.columns:
        y = df['WeightCategory']
        X = df.drop(columns=['WeightCategory'])
    else:
        X = df.copy()

    #Applying onehot encoders and standard scalers on
    categorical_features = ['Gender', 'MTRANS', 'FAVC', 'SCC', 'family_history_w'
    ordinal_features_caec = ['CAEC']
    ordinal_features_calc = ['CALC']
    numerical_features = ['Age', 'Height', 'Weight', 'HealthyEatingScore_PCA']
    remainder_features = ['FCVC', 'NCP', 'CH20', 'FAF', 'TUE']

    onehot_encoder = OneHotEncoder(drop='first', sparse_output=False)
    X_categorical = onehot_encoder.fit_transform(X[categorical_features])
    categorical_feature_names = onehot_encoder.get_feature_names_out(categorical_
    
    ordinal_encoder_caec = OrdinalEncoder(categories=[['no', 'Sometimes', 'Frequ
    X_ordinal_caec = ordinal_encoder_caec.fit_transform(X[ordinal_features_caec])
    ordinal_caec_feature_names = ordinal_encoder_caec.get_feature_names_out(ordin

    ordinal_encoder_calc = OrdinalEncoder(categories=[['no', 'Sometimes', 'Frequ
    X_ordinal_calc = ordinal_encoder_calc.fit_transform(X[ordinal_features_calc])
    ordinal_calc_feature_names = ordinal_encoder_calc.get_feature_names_out(ordin

    scaler = StandardScaler()
    X_numerical = scaler.fit_transform(X[numerical_features])
    numerical_feature_names = scaler.get_feature_names_out(numerical_features)

    X_remainder = X[remainder_features].values
    remainder_feature_names = remainder_features

    X_transformed = np.concatenate([X_categorical, X_ordinal_caec, X_ordinal_cal
    all_feature_names = np.concatenate([
        categorical_feature_names,
        ordinal_caec_feature_names,
        ordinal_calc_feature_names,
        numerical_feature_names,
        remainder_feature_names
```

```
])
X = pd.DataFrame(X_transformed, columns=all_feature_names, index=X.index)

return X, y
```

This function is basically transforming our data into X and y using scaling, encoding and adding the health score column. We made it into a function so that it could be reused on the training and the test dataset.

```
X, y = preprocess_data(df)
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.25,random_state=42)
print(X_train.shape,X_test.shape,y_train.shape,y_test.shape)

(11649, 20) (3884, 20) (11649,) (3884,)
```

▼ 5. Testing Different models

```
dt = DecisionTreeClassifier()
ada = AdaBoostClassifier()
rf = RandomForestClassifier()
xgb = XGBClassifier()
```

```
model_dt = dt.fit(X_train,y_train)
model_ada = ada.fit(X_train,y_train)
model_rf = rf.fit(X_train,y_train)
model_xgb = xgb.fit(X_train,y_train)

y_pred_dt = model_dt.predict(X_test)
y_pred_ada = model_ada.predict(X_test)
y_pred_rf = model_rf.predict(X_test)
y_pred_xgb = model_xgb.predict(X_test)

print(accuracy_score(y_test,y_pred_dt))
print(accuracy_score(y_test,y_pred_ada))
print(accuracy_score(y_test,y_pred_rf))
print(accuracy_score(y_test,y_pred_xgb))
```

```
0.8493820803295572
0.6967044284243048
0.8954685890834192
0.9034500514933058
```

Here we run different classification models on the dataset and we see that XG Boost performs the best here.

```
parameters = {'gamma': [1.0935154163692284], 'learning_rate':[0.01640914713624878  
    'max_depth':[7], 'max_leaves': [31],  
    'min_child_weight':[5], 'n_estimators': [1032], 'reg_lambda': [2.25]  
    'reg_alpha': [0.02591342917807464], 'subsample': [0.71045525886008]  
    'colsample_bytree': [0.5719600338322923]  
}
```

These are the values we passed to GridsearchCV to test our model on. It had a lot of values but through a lot of trial and error we reached to these values in the end. We removed the other values as running this cell took a lot of time as the number of values was high.

```
xgb_tune = XGBClassifier(n_jobs=-1)  
grid_search = GridSearchCV(estimator=xgb_tune, param_grid=parameters, scoring='accuracy')  
grid_search.fit(X_train, y_train)  
print(grid_search.best_score_)  
print(grid_search.best_params_)  
best_model = grid_search.best_estimator_
```

0.903081799368662
{'colsample_bytree': 0.5719600338322923, 'gamma': 1.0935154163692284, 'learning_rate': 0.01640914713624878}

```
y_pred = best_model.predict(X_test)  
accuracy_score(y_test, y_pred)
```

0.9111740473738414

```
cm = confusion_matrix(y_test, y_pred)  
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=list(rev_map.values()), yticklabels=list(rev_map.values()))  
plt.xlabel('Predicted Weight Category')  
plt.ylabel('Actual Weight Category')  
plt.title('Confusion Matrix for Best Model')  
plt.show()
```

Confusion Matrix for Best Model